

NPTTEL | DEEP LEARNING | CS1075

Reticular theory is an obsolete scientific theory in neurobiology that stated that everything in the nervous system, such as brain, is a single continuous network. The concept was postulated by a German anatomist Joseph von Gerlach in 1871, and was most popularised by the Nobel laureate Italian physician Camillo Golgi.

Reticular theory – single cell; Joseph von Gerlach, 1871

Camillo Golgi – proponent of Reticular theory (single continuous network); Staining technique

Santiago Cajal – not a single one; collection of neurons – neuron doctrine!

Heinrich Gottfried – coined neuron and chromosome; consolidated into Neuron Doctrine!

1906 Nobel in Medicine – Camillo Golgi, Santiago Ramón y Cajal

Electron microscope – gap between neurons; 1950's, debate settled!

McCulloch-Pitts Neuron – 1943 + AI

AI – 1956, coined the word AI

Perceptron –

Deep Learning – large number of neurons, inter-connected to perform certain activities

Limitations – it's possible that a perceptron cannot make simple decisions; not able to do basic XOR.

Types of AI – Connectionist AI and Symbolic AI; we'll focus on Connectionism

"A multi-layered n/w of neurons can do – whatever cited as limitation {a single neuron cannot do simple tasks!}"

Back Propagation

Gradient descent

Universal Approximation theorem

McCulloch Pitts neuron [MP neuron]

Threshold –theta

Exhibitors and inhibitors

Classifying 0 or 1 / linear decision boundary / linearly separable function

Multi-dimensional planes

Bias ($w_0 = -\text{theta}$)

Sigma $i=0$ to n , $w_0, w_1, w_2, \dots, w_n$

Perceptron

Give data

M movies and labels (class) – liked or not/ binary decision

Data+labels

Perceptron is supposed to Adjust the weights – in such a way that we should be able to separate

Perfect match

Feed movies – return same label

PLA – Perceptron Learning Algorithm

P -> inputs with label 1

N -> inputs with label 0

Assign weights randomly!

As we don't know the weights as of now....

Convergence – when there're no more errors...

$W = [w_0, w_1, w_2, \dots, w_n] / n+1$ dimension

X is n -dimension

$X = [1, x_1, x_2, \dots, x_n]$

While convergence != "do

Pick random x element of P Union N

If X element of P AND If summation ($w^T X$) > 0 THEN

$$W = W + X$$

If X element of N AND If summation ($w^T X$) ≤ 0

$$W = W - X$$

Dot Product:

$$w \cdot x = w^T X = \text{SUM}(w_i * x_i) ; i = 0 \text{ to } n$$

Perceptron rule:

$$y = 1 \text{ if } w^T X \geq 0$$

$$y = 0 \text{ if } w^T X < 0$$


$w^T X \Rightarrow$ divides the input space into two...

Need the line $w^T X = 0$; perpendicular

Every point (x) on the line satisfies the equation $w^T X = 0$


Angle $\alpha = 90^\circ / \cos 90 = 0$

Consider some points (vectors) which lie in the positive half space of this line (i.e., $w^T x \geq 0$)



0)

- What will be the angle between any such vector and \mathbf{w} ? Obviously, less than 90°
- What about points (vectors) which lie in the negative half space of this line (i.e., $\mathbf{w}^T \mathbf{x} < 0$)
- What will be the angle between any such vector and \mathbf{w} ? Obviously, greater than 90°



- Of course, this also follows from the formula

$$(\cos \alpha = \frac{\mathbf{w}^T \mathbf{x}}{\|\mathbf{w}\| \|\mathbf{x}\|})$$

- For $\mathbf{x} \in P$ if $\mathbf{w} \cdot \mathbf{x} < 0$ then it means that the angle (α) between this \mathbf{x} and the current \mathbf{w} is greater than 90° (but we want α to be less than 90°)
- What happens to the new angle (α_{new}) when $\mathbf{w}_{new} = \mathbf{w} + \mathbf{x}$

$$\begin{aligned} \cos(\alpha_{new}) &\propto \mathbf{w}_{new}^T \mathbf{x} \\ &\propto (\mathbf{w} + \mathbf{x})^T \mathbf{x} \\ &\propto \mathbf{w}^T \mathbf{x} + \mathbf{x}^T \mathbf{x} \\ &\propto \cos \alpha + \mathbf{x}^T \mathbf{x} \\ \cos(\alpha_{new}) &> \cos \alpha \end{aligned}$$

- For $\mathbf{x} \in N$ if $\mathbf{w} \cdot \mathbf{x} \geq 0$ then it means that the angle (α) between this \mathbf{x} and the current \mathbf{w} is less than 90° (but we want α to be greater than 90°)
- What happens to the new angle (α_{new}) when $\mathbf{w}_{new} = \mathbf{w} - \mathbf{x}$

$$\begin{aligned} \cos(\alpha_{new}) &\propto \mathbf{w}_{new}^T \mathbf{x} \\ &\propto (\mathbf{w} - \mathbf{x})^T \mathbf{x} \\ &\propto \mathbf{w}^T \mathbf{x} - \mathbf{x}^T \mathbf{x} \\ &\propto \cos \alpha - \mathbf{x}^T \mathbf{x} \end{aligned}$$

$$\cos(\alpha_{new}) < \cos \alpha$$

- Thus α_{new} will be greater than α and this is exactly what we want

McCulloch Pitts Neuron

(assuming no inhibitory inputs)

$$y = 1 \quad \text{if } \sum_{i=0}^n x_i \geq \theta$$


$$= 0 \quad \text{if } \sum_{i=0}^n x_i < \theta$$

Perceptron

$$y = 1 \quad \text{if } \sum_{i=0}^n w_i * x_i \geq \theta$$

$$= 0 \quad \text{if } \sum_{i=0}^n w_i * x_i < \theta$$

- From the equations it should be clear that even a perceptron separates the input space into two halves
- All inputs which produce a 1 lie on one side and all inputs which produce a 0 lie on the other side
- In other words, a single perceptron can only be used to implement linearly separable functions
- Then what is the difference? The weights (including threshold) can be learned and the inputs can be real valued



c. **Apply** that weighted sum to the correct Activation Function.

For Example : Unit Step Activation Function.




Fig: Unit Step Activation Function

Why do we need Activation Function?

In short, the activation functions are used to map the input between the required values like (0, 1) or (-1, 1).

a. All the inputs x are multiplied with their weights w . Let's call it k .




Fig: Multiplying inputs with weights for 5 inputs

b. **Add** all the multiplied values and call them *Weighted Sum*.


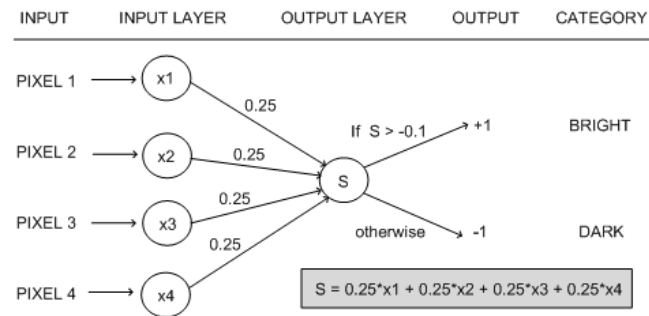


Fig: Adding with Summation

Why do we need Weights and Bias?


Weights shows the strength of the particular node.


A bias value allows you to shift the activation function curve up or down.




Where we use Perceptron?

Perceptron is usually used to classify the data into two parts. Therefore, it is also known as a Linear Binary Classifier.





- We initialized \mathbf{w} to a random value
- We observe that currently, $\mathbf{w} \cdot \mathbf{x} < 0$ (\because angle $> 90^\circ$) for all the positive points and $\mathbf{w} \cdot \mathbf{x} \geq 0$ (\because angle $< 90^\circ$) for all the negative points (the situation is exactly oppsite of what we actually want it to be)
- We now run the algorithm by randomly going over the points
- Randomly pick a point (say, p_2), apply correction $\mathbf{w} = \mathbf{w} + \mathbf{x}$ \because $\mathbf{w} \cdot \mathbf{x} < 0$ (you can check the angle visually)



- We initialized \mathbf{w} to a random value
- We observe that currently, $\mathbf{w} \cdot \mathbf{x} < 0$ (\because angle $> 90^\circ$) for all the positive points and $\mathbf{w} \cdot \mathbf{x} \geq 0$ (\because angle $< 90^\circ$) for all the negative points (the situation is exactly oppsite of what we actually want it to be)
- We now run the algorithm by randomly going over the points
- Randomly pick a point (say, n_1), no correction needed \because $\mathbf{w} \cdot \mathbf{x} < 0$ (you can check the angle visually)

Theorem

Definition: Two sets P and N of points in an n -dimensional space are called absolutely linearly separable if $n + 1$ real numbers w_0, w_1, \dots, w_n exist such that every point $(x_1, x_2, \dots, x_n) \in P$ satisfies $\sum_{i=1}^n w_i * x_i > w_0$ and every point $(x_1, x_2, \dots, x_n) \in N$ satisfies $\sum_{i=1}^n w_i * x_i < w_0$.

W2

*function which
is linearly separable
OR*

*0 0 0
0 1 1
1 0 1
1 1 1*

*is not linearly separable
XOR*

*0 0 0
0 1 1
1 0 1
1 1 0*

*THERE DOESN'T EXIST A LINE
WHICH SEPARATES X-Y*

MOST REAL WORLD DATA IS NOT LINEARLY SEPARABLE

*n inputs
 2^{2n} functions
- how many are linearly separable?*


*2 inputs
16 functions
2 are not linearly separable
XOR and !XOR*

*HOW any BOOLEAN function CAN be REPRESENTED using A perceptron
- model/ input/ truth table/ output/*

Boolean function!


*REPRESENTED – feeding any value of x,y
- Network will give the same output as that of the truth table
-*

0 or -1




- For this discussion, we will assume True = +1 and False = -1
- We consider 2 inputs and 4 perceptrons
- Each input is connected to all the 4 perceptrons with specific weights
- The bias (w_0) of each perceptron is -2 (i.e., each perceptron will fire only if the weighted sum of its input is ≥ 2)
- Each of these perceptrons is connected to an output perceptron by weights (which need to be learned)
- The output of this perceptron (y) is the output of this network

- This network contains 3 layers
- The layer containing the inputs (x_1, x_2) is called the **input layer**
- The middle layer containing the 4 perceptrons is called the **hidden layer**
- The final layer containing one output neuron is called the **output layer**
- The outputs of the 4 perceptrons in the hidden layer are denoted by h_1, h_2, h_3, h_4



- For this discussion, we will assume True = +1 and False = -1
- We consider 2 inputs and 4 perceptrons
- Each input is connected to all the 4 perceptrons with specific weights
- The bias (w_0) of each perceptron is -2 (i.e., each perceptron will fire only if the weighted sum of its input is ≥ 2)
- Each of these perceptrons is connected to an output perceptron by weights (which need to be learned)
- The output of this perceptron (y) is the output of this network


Terminology:



- This network contains 3 layers
- The layer containing the inputs (x_1, x_2) is called the **input layer**
- The middle layer containing the 4 perceptrons is called the **hidden layer**
- The final layer containing one output neuron is called the **output layer**
- The outputs of the 4 perceptrons in the hidden layer are denoted by h_1, h_2, h_3, h_4


- We claim that this network can be used to implement **any** boolean function (linearly separable or not) !
- In other words, we can find w_1, w_2, w_3, w_4 such that the truth table of any boolean function can be represented by this network
- Astonishing claim! Well, not really, if you understand what is going on

- Each perceptron in the middle layer fires only for a specific input (and no two perceptrons fire for the same input)
- Let us see why this network works by taking an example of the XOR function



red edge indicates $w = -1$
blue edge indicates $w = +1$

- implement **any** boolean function (linearly separable or not) !
- In other words, we can find w_1, w_2, w_3, w_4 such that the truth table of any boolean function can be represented by this network
- Astonishing claim! Well, not really, if you understand what is going on
- Each perceptron in the middle layer fires only for a specific input (and no two perceptrons fire for the same input)
- the fourth perceptron fires for $\{1,1\}$




red edge indicates $w = -1$
blue edge indicates $w = +1$

- Let w_0 be the bias output of the neuron (i.e., it will fire if $\sum_{i=1}^4 w_i h_i \geq w_0$)

x_1	x_2	XOR	h_1	h_2	h_3	h_4	$\sum_{i=1}^4 w_i h_i$
0	0	0	1	0	0	0	w_1
0	1	1	0	1	0	0	w_2
1	0	1	0	0	1	0	w_3
1	1	0	0	0	0	1	w_4

- This results in the following four conditions to implement XOR: $w_1 < w_0, w_2 \geq w_0, w_3 \geq w_0, w_4 < w_0$



Theorem

Any boolean function of n inputs can be represented exactly by a network of perceptrons containing 1 hidden layer with 2^n perceptrons and one output layer containing 1 perceptron

Proof (informal:) We just saw how to construct such a network

Note: A network of $2^n + 1$ perceptrons is not necessary but sufficient. For example, we already saw how to represent AND function with just 1 perceptron

Catch: As n increases the number of perceptrons in the hidden layers obviously increases exponentially

TAKEAWAY:

PERCEPTRON CAN IMPLEMENT ANY BOOLEAN FUNCTION, WHETHER LINEARLY SEPARABLE OR NOT

SIGMOID FUNCTION

OIL wells

Function which takes x (vector with n variables) and gives y

- Whether to drill or not


-

The story ahead ...


- Enough about boolean functions!
- What about arbitrary functions of the form $y = f(x)$ where $x \in \mathbb{R}^n$ (instead of $\{0, 1\}^n$) and $y \in \mathbb{R}$ (instead of $\{0, 1\}$) ?

Handwritten expression: $x \begin{bmatrix} \mathbb{R} & \mathbb{R} \\ x_1 & x_2 & x_3 & \dots & x_n \end{bmatrix} \in \mathbb{R}^n$. Below the vector, the words "Salinity Density Pressure" are written, suggesting these as possible components of the vector x .

- From arbitrary neurons to sigmoid neurons
 - Real functions
 - Perceptron will fire when the weighted sum is greater than the threshold
- Networks which could represent Boolean functions
 - 1 hidden layer minimum
 - Hidden layer grows exponentially!
 - Harsh logic used by perceptron
 - 0.49 and 0.51 lies pretty close but treated differently by perceptron functions
 -




- The thresholding logic used by a perceptron is very harsh !
- For example, let us return to our problem of deciding whether we will like or dislike a movie
- Consider that we base our decision only on one input ($x_1 = \text{criticsRating}$ which lies between 0 and 1)
- If the threshold is 0.5 ($w_0 = -0.5$) and $w_1 = 1$ then what would be the decision for a movie with $\text{criticsRating} = 0.51$?



- This behavior is not a characteristic of the specific problem we chose or the specific weight and threshold that we chose
- It is a characteristic of the perceptron function itself which behaves like a step function

Sigmoid Neuron



- Introducing sigmoid neurons where the output function is much smoother than the step function
- Here is one form of the sigmoid function called the logistic function

$$y = \frac{1}{1 + e^{-(w_0 + \sum_{i=1}^n w_i x_i)}}$$

$w^T x = \sum_{i=0}^n w_i x_i$
 $\frac{1}{1 + e^{-w^T x}}$
 $w^T x \rightarrow \infty$

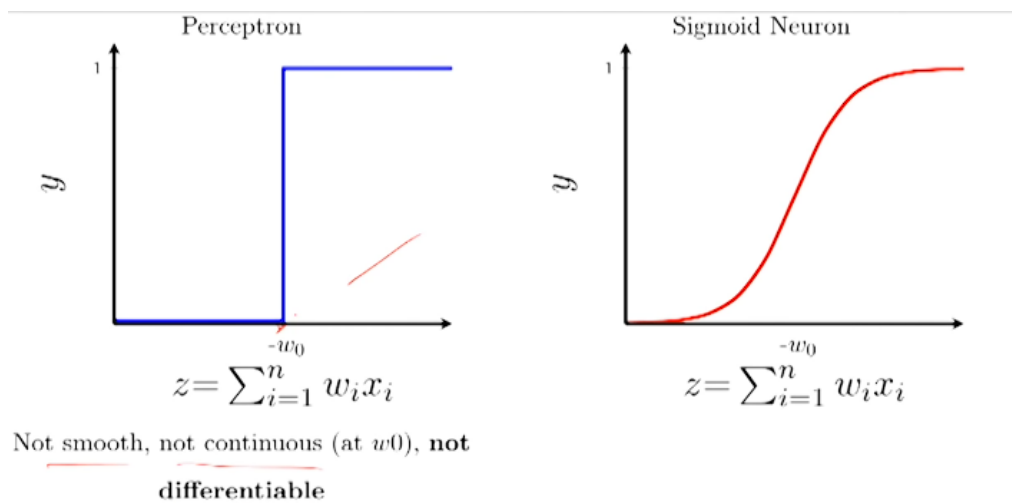
$$\frac{1}{1 + e^{-w^T x}} \quad w^T x = 0$$

$w^T x \rightarrow \infty$
 $w^T x \rightarrow -\infty$



Sigmoid function range lies between 0 and 1

Same goes with probability

Not harsh anymore!



Sigmoid (logistic) Neuron



- Earlier we mentioned that a single perceptron cannot deal with this data because it is not linearly separable
- What does “cannot deal with” mean?
- What would happen if we use a perceptron model to classify this data ?
- We would probably end up with a line like this ...
- This line doesn't seem to be too bad

This brings us to a typical machine learning setup which has the following components...

- **Data:** $\{x_i, y_i\}_{i=1}^n$
- **Model:** Our approximation of the relation between \mathbf{x} and y . For example,

$$\hat{y} = \frac{1}{1 + e^{-(\mathbf{w}^T \mathbf{x})}}$$

or $\hat{y} = \mathbf{w}^T \mathbf{x}$

or $\hat{y} = \mathbf{x}^T \mathbf{W} \mathbf{x}$

for non-linear

or just about any function

- **Parameters:** In all the above cases, w is a parameter which needs to be learned from the data

Learning algorithm – algorithm for learning the parameters (w) of the model!

- Gradient descent
- Perceptron learning algorithm

As an illustration, consider our movie example

- **Data:** $\{x_i = \text{movie}, y_i = \text{like/dislike}\}_{i=1}^n$
- **Model:** Our approximation of the relation between \mathbf{x} and y (the probability of liking a movie).

$$\hat{y} = \frac{1}{1 + e^{-(\mathbf{w}^T \mathbf{x})}}$$

- **Parameter:** w
- **Learning algorithm:** Gradient Descent [we will see soon]
- **Objective/Loss/Error function:** One possibility is

$$\mathcal{L}(\mathbf{w}) = \sum_{i=1}^n (\hat{y}_i - y_i)^2$$



The learning algorithm should aim to find a w which function (squared error between y and \hat{y})

« »